

```
In [ ]: !pip install scikit-learn
```

```
In [11]: pip install matplotlib
```

Requirement already satisfied: matplotlib in c:\users\anisha\anaconda3\lib\site-packages (3.4.3)  
 Requirement already satisfied: pyparsing>=2.2.1 in c:\users\anisha\anaconda3\lib\site-packages (from matplotlib) (3.0.4)  
 Requirement already satisfied: python-dateutil>=2.7 in c:\users\anisha\anaconda3\lib\site-packages (from matplotlib) (2.8.2)  
 Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\anisha\anaconda3\lib\site-packages (from matplotlib) (1.3.1)  
 Requirement already satisfied: numpy>=1.16 in c:\users\anisha\anaconda3\lib\site-packages (from matplotlib) (1.20.3)  
 Requirement already satisfied: pillow>=6.2.0 in c:\users\anisha\anaconda3\lib\site-packages (from matplotlib) (9.2.0)  
 Requirement already satisfied: cycler>=0.10 in c:\users\anisha\anaconda3\lib\site-packages (from matplotlib) (0.10.0)  
 Requirement already satisfied: six in c:\users\anisha\anaconda3\lib\site-packages (from cycler>=0.10->matplotlib) (1.16.0)  
 Note: you may need to restart the kernel to use updated packages.

```
In [1]: import numpy as np
        from sklearn.linear_model import LinearRegression
```

## Multiple Linear Regression on startup dataset

```
In [3]: !pip install pandas
```

Requirement already satisfied: pandas in c:\users\india\anaconda3\lib\site-packages (0.25.1)  
 Requirement already satisfied: pytz>=2017.2 in c:\users\india\anaconda3\lib\site-packages (from pandas) (2019.3)  
 Requirement already satisfied: python-dateutil>=2.6.1 in c:\users\india\anaconda3\lib\site-packages (from pandas) (2.8.0)  
 Requirement already satisfied: numpy>=1.13.3 in c:\users\india\anaconda3\lib\site-packages (from pandas) (1.16.5)  
 Requirement already satisfied: six>=1.5 in c:\users\india\anaconda3\lib\site-packages (from python-dateutil>=2.6.1->pandas) (1.12.0)

```
In [1]: import pandas as pd
        df = pd.read_csv(r"C:\Users\Anisha\Downloads\50_Startups.csv")
        df
```

```
Out[1]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

	R&D Spend	Administration	Marketing Spend	State	Profit
5	131876.90	99814.71	362861.36	New York	156991.12
6	134615.46	147198.87	127716.82	California	156122.51
7	130298.13	145530.06	323876.68	Florida	155752.60
8	120542.52	148718.95	311613.29	New York	152211.77
9	123334.88	108679.17	304981.62	California	149759.96
10	101913.08	110594.11	229160.95	Florida	146121.95
11	100671.96	91790.61	249744.55	California	144259.40
12	93863.75	127320.38	249839.44	Florida	141585.52
13	91992.39	135495.07	252664.93	California	134307.35
14	119943.24	156547.42	256512.92	Florida	132602.65
15	114523.61	122616.84	261776.23	New York	129917.04
16	78013.11	121597.55	264346.06	California	126992.93
17	94657.16	145077.58	282574.31	New York	125370.37
18	91749.16	114175.79	294919.57	Florida	124266.90
19	86419.70	153514.11	0.00	New York	122776.86
20	76253.86	113867.30	298664.47	California	118474.03
21	78389.47	153773.43	299737.29	New York	111313.02
22	73994.56	122782.75	303319.26	Florida	110352.25
23	67532.53	105751.03	304768.73	Florida	108733.99
24	77044.01	99281.34	140574.81	New York	108552.04
25	64664.71	139553.16	137962.62	California	107404.34
26	75328.87	144135.98	134050.07	Florida	105733.54
27	72107.60	127864.55	353183.81	New York	105008.31
28	66051.52	182645.56	118148.20	Florida	103282.38
29	65605.48	153032.06	107138.38	New York	101004.64
30	61994.48	115641.28	91131.24	Florida	99937.59
31	61136.38	152701.92	88218.23	New York	97483.56
32	63408.86	129219.61	46085.25	California	97427.84
33	55493.95	103057.49	214634.81	Florida	96778.92
34	46426.07	157693.92	210797.67	California	96712.80
35	46014.02	85047.44	205517.64	New York	96479.51
36	28663.76	127056.21	201126.82	Florida	90708.19
37	44069.95	51283.14	197029.42	California	89949.14
38	20229.59	65947.93	185265.10	New York	81229.06
39	38558.51	82982.09	174999.30	California	81005.76
40	28754.33	118546.05	172795.67	California	78239.91

	R&D Spend	Administration	Marketing Spend	State	Profit
41	27892.92	84710.77	164470.71	Florida	77798.83
42	23640.93	96189.63	148001.11	California	71498.49
43	15505.73	127382.30	35534.17	New York	69758.98
44	22177.74	154806.14	28334.72	California	65200.33
45	1000.23	124153.04	1903.93	New York	64926.08
46	1315.46	115816.21	297114.46	Florida	49490.75
47	0.00	135426.92	0.00	California	42559.73
48	542.05	51743.15	0.00	New York	35673.41
49	0.00	116983.80	45173.06	California	14681.40

In [2]:

```
shape=df.shape
print("Dataset contains {} rows and {} columns".format(shape[0],shape[1]))
```

Dataset contains 50 rows and 5 columns

In [3]:

```
df.columns
```

Out[3]:

```
Index(['R&D Spend', 'Administration', 'Marketing Spend', 'State', 'Profit'], dtype='object')
```

In [4]:

```
df.describe()
```

Out[4]:

	R&D Spend	Administration	Marketing Spend	Profit
<b>count</b>	50.000000	50.000000	50.000000	50.000000
<b>mean</b>	73721.615600	121344.639600	211025.097800	112012.639200
<b>std</b>	45902.256482	28017.802755	122290.310726	40306.180338
<b>min</b>	0.000000	51283.140000	0.000000	14681.400000
<b>25%</b>	39936.370000	103730.875000	129300.132500	90138.902500
<b>50%</b>	73051.080000	122699.795000	212716.240000	107978.190000
<b>75%</b>	101602.800000	144842.180000	299469.085000	139765.977500
<b>max</b>	165349.200000	182645.560000	471784.100000	192261.830000

In [5]:

```
c = df.corr()
c
#Inference: We can see that all three columns have a direct relationship with the pr
```


Out[5]:

	R&D Spend	Administration	Marketing Spend	Profit
<b>R&amp;D Spend</b>	1.000000	0.241955	0.724248	0.972900
<b>Administration</b>	0.241955	1.000000	-0.032154	0.200717
<b>Marketing Spend</b>	0.724248	-0.032154	1.000000	0.747766
<b>Profit</b>	0.972900	0.200717	0.747766	1.000000

## Correlation Matrix

```
In [13]: import seaborn as sns # for visualization
import matplotlib as plt
sns.heatmap(c,annot=True,cmap='Blues')
plt.show()
```

Out[13]: <AxesSubplot:>



```
In [14]: # splitting Dataset in Dependent & Independent Variables
X = df.iloc[:, :-1].values
y = df.iloc[:, 4].values
```

In [15]: X

```
Out[15]: array([[165349.2, 136897.8, 471784.1, 'New York'],
 [162597.7, 151377.59, 443898.53, 'California'],
 [153441.51, 101145.55, 407934.54, 'Florida'],
 [144372.41, 118671.85, 383199.62, 'New York'],
 [142107.34, 91391.77, 366168.42, 'Florida'],
 [131876.9, 99814.71, 362861.36, 'New York'],
 [134615.46, 147198.87, 127716.82, 'California'],
 [130298.13, 145530.06, 323876.68, 'Florida'],
 [120542.52, 148718.95, 311613.29, 'New York'],
 [123334.88, 108679.17, 304981.62, 'California'],
 [101913.08, 110594.11, 229160.95, 'Florida'],
 [100671.96, 91790.61, 249744.55, 'California'],
 [93863.75, 127320.38, 249839.44, 'Florida'],
 [91992.39, 135495.07, 252664.93, 'California'],
 [119943.24, 156547.42, 256512.92, 'Florida'],
 [114523.61, 122616.84, 261776.23, 'New York'],
 [78013.11, 121597.55, 264346.06, 'California'],
 [94657.16, 145077.58, 282574.31, 'New York'],
 [91749.16, 114175.79, 294919.57, 'Florida'],
 [86419.7, 153514.11, 0.0, 'New York'],
 [76253.86, 113867.3, 298664.47, 'California'],
 [78389.47, 153773.43, 299737.29, 'New York'],
 [73994.56, 122782.75, 303319.26, 'Florida'],
 [67532.53, 105751.03, 304768.73, 'Florida'],
```

```
[77044.01, 99281.34, 140574.81, 'New York'],
[64664.71, 139553.16, 137962.62, 'California'],
[75328.87, 144135.98, 134050.07, 'Florida'],
[72107.6, 127864.55, 353183.81, 'New York'],
[66051.52, 182645.56, 118148.2, 'Florida'],
[65605.48, 153032.06, 107138.38, 'New York'],
[61994.48, 115641.28, 91131.24, 'Florida'],
[61136.38, 152701.92, 88218.23, 'New York'],
[63408.86, 129219.61, 46085.25, 'California'],
[55493.95, 103057.49, 214634.81, 'Florida'],
[46426.07, 157693.92, 210797.67, 'California'],
[46014.02, 85047.44, 205517.64, 'New York'],
[28663.76, 127056.21, 201126.82, 'Florida'],
[44069.95, 51283.14, 197029.42, 'California'],
[20229.59, 65947.93, 185265.1, 'New York'],
[38558.51, 82982.09, 174999.3, 'California'],
[28754.33, 118546.05, 172795.67, 'California'],
[27892.92, 84710.77, 164470.71, 'Florida'],
[23640.93, 96189.63, 148001.11, 'California'],
[15505.73, 127382.3, 35534.17, 'New York'],
[22177.74, 154806.14, 28334.72, 'California'],
[1000.23, 124153.04, 1903.93, 'New York'],
[1315.46, 115816.21, 297114.46, 'Florida'],
[0.0, 135426.92, 0.0, 'California'],
[542.05, 51743.15, 0.0, 'New York'],
[0.0, 116983.8, 45173.06, 'California']], dtype=object)
```

```
In [16]: from sklearn.preprocessing import LabelEncoder
```

```
In [17]: labelencoder = LabelEncoder()
X[:, 3] = labelencoder.fit_transform(X[:, 3])
X1 = pd.DataFrame(X)
X1.head()
```

```
Out[17]:
```

	0	1	2	3
0	165349.2	136897.8	471784.1	2
1	162597.7	151377.59	443898.53	0
2	153441.51	101145.55	407934.54	1
3	144372.41	118671.85	383199.62	2
4	142107.34	91391.77	366168.42	1

```
In [18]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test = train_test_split(X,y,train_size=0.7,random_state=0)
x_train
```

```
Out[18]: array([[130298.13, 145530.06, 323876.68, 1],
 [119943.24, 156547.42, 256512.92, 1],
 [1000.23, 124153.04, 1903.93, 2],
 [542.05, 51743.15, 0.0, 2],
 [65605.48, 153032.06, 107138.38, 2],
 [114523.61, 122616.84, 261776.23, 2],
 [61994.48, 115641.28, 91131.24, 1],
 [63408.86, 129219.61, 46085.25, 0],
 [78013.11, 121597.55, 264346.06, 0],
 [23640.93, 96189.63, 148001.11, 0],
```

```
[76253.86, 113867.3, 298664.47, 0],
[15505.73, 127382.3, 35534.17, 2],
[120542.52, 148718.95, 311613.29, 2],
[91992.39, 135495.07, 252664.93, 0],
[64664.71, 139553.16, 137962.62, 0],
[131876.9, 99814.71, 362861.36, 2],
[94657.16, 145077.58, 282574.31, 2],
[28754.33, 118546.05, 172795.67, 0],
[0.0, 116983.8, 45173.06, 0],
[162597.7, 151377.59, 443898.53, 0],
[93863.75, 127320.38, 249839.44, 1],
[44069.95, 51283.14, 197029.42, 0],
[77044.01, 99281.34, 140574.81, 2],
[134615.46, 147198.87, 127716.82, 0],
[67532.53, 105751.03, 304768.73, 1],
[28663.76, 127056.21, 201126.82, 1],
[78389.47, 153773.43, 299737.29, 2],
[86419.7, 153514.11, 0.0, 2],
[123334.88, 108679.17, 304981.62, 0],
[38558.51, 82982.09, 174999.3, 0],
[1315.46, 115816.21, 297114.46, 1],
[144372.41, 118671.85, 383199.62, 2],
[165349.2, 136897.8, 471784.1, 2],
[0.0, 135426.92, 0.0, 0],
[22177.74, 154806.14, 28334.72, 0]], dtype=object)
```

In [19]: `from sklearn.linear_model import LinearRegression`

```
model = LinearRegression()
model.fit(x_train,y_train)
print('Model has been trained successfully')
```

Model has been trained successfully

In [20]: `y_pred = model.predict(x_test)`  
`y_pred`

Out[20]: `array([104055.1842384 , 132557.60289702, 133633.01284474, 72336.28081054,`  
`179658.27210893, 114689.63133397, 66514.82249033, 98461.69321326,`  
`114294.70487032, 169090.51127461, 96281.907934 , 88108.30057881,`  
`110687.1172322 , 90536.34203081, 127785.3793861 ])`

In [21]: `testing_data_model_score = model.score(x_test, y_test)`  
`print("Model Score/Performance on Testing data",testing_data_model_score)`  
  
`training_data_model_score = model.score(x_train, y_train)`  
`print("Model Score/Performance on Training data",training_data_model_score)`

Model Score/Performance on Testing data 0.9355139722149947  
 Model Score/Performance on Training data 0.9515496105627431

In [27]: `# Compare predicted and actual values`  
`df1 = pd.DataFrame(data={'Predicted value':y_pred.flatten(),'Actual Value':y_test.fl`  
`df1`

Out[27]:

	Predicted value	Actual Value
0	104055.184238	103282.38
1	132557.602897	144259.40

	Predicted value	Actual Value
2	133633.012845	146121.95
3	72336.280811	77798.83
4	179658.272109	191050.39
5	114689.631334	105008.31
6	66514.822490	81229.06
7	98461.693213	97483.56
8	114294.704870	110352.25
9	169090.511275	166187.94
10	96281.907934	96778.92
11	88108.300579	96479.51
12	110687.117232	105733.54
13	90536.342031	96712.80
14	127785.379386	124266.90

### Model evaluation

R2 score: R2 score – R squared score. It is one of the statistical approaches by which we can find the variance or the spread of the target and feature data.

```
In [22]: from sklearn.metrics import r2_score

r2Score = r2_score(y_pred, y_test)
print("R2 score of model is :", r2Score*100)
```

R2 score of model is : 93.39448007716634

MSE: MSE – Mean Squared Error. By using this approach we can find that how much the regression best fit line is close to all the residual.

```
In [29]: from sklearn.metrics import mean_squared_error

mse = mean_squared_error(y_pred, y_test)
print("Mean Squarred Error is :", mse*100)
```

Mean Squarred Error is : 6224496238.946446

RMSE: RMSE – Root Mean Squared Error. This is similar to the Mean squared error(MSE) approach, the only difference is that here we find the root of the mean squared error i.e. root of the Mean squared error is equal to Root Mean Squared Error. The reason behind finding the root is to find the more close residual to the values found by mean squared error.

```
In [30]: import numpy as np

rmse = np.sqrt(mean_squared_error(y_pred, y_test))
print("Root Mean Squarred Error is :", rmse*100)
```

Root Mean Squarred Error is : 788954.7666974607

MAE: MAE – Mean Absolute Error. By using this approach we can find the difference between

the actual values and predicted values but that difference is absolute i.e. the difference is positive.

```
In [31]: from sklearn.metrics import mean_absolute_error

mae = mean_absolute_error(y_pred,y_test)
print("Mean Absolute Error is : " ,mae)
```

Mean Absolute Error is : 6503.57732358002

In [ ]: